

09.3.- Data Cleansing-Outliers_CU_18_20_infra_meteo_v_01

June 13, 2023

#

CU18_Infraestructuras_eventos

Citizenlab Data Science Methodology > II - Data Processing Domain *** > # 09.3.- Data Cleansing
- Outliers

Data Cleaning refers to identifying and correcting (or removing) errors in the dataset that may negatively impact a predictive model, replacing, modifying, or deleting the dirty or coarse data.

0.1 Tasks

Basic operations	Text data analysis
	Delete Needless/Irrelevant/Private Columns Inconsistent Data. Expected values Zeroes Columns with a Single Value Columns with Very Few Values Columns with Low Variance Duplicates (rows/samples) & (columns/features) Data
Missing Values	Missing Values Identification
	Missing Values Per Sample Missing Values Per Feature Zero Missing Values Other Missing Values Null/NaN Missing Values Delete Missing Values Deleting Rows with Missing Values in Target Column Deleting Rows with Missing Values Deleting Features with some Missing Values Deleting Features using Rate Missing Values
	Basic Imputation
	Imputation by Previous Row Value Imputation by Next Row Value
	Statistical Imputation
	Selection of Imputation Strategy Constant Imputation Mean Imputation Median Imputation Most Frequent Imputation Interpolation Imputation
	Prediction Imputation (KNN Imputation)
	Evaluating k-hyperparameter in KNN Imputation Applying KNN Imputation
	Iterative Imputation
	Evaluating Different Imputation Order Applying Iterative Imputation
Outliers	Outliers - Univariate
	Visualizing Outliers Distribution Box Plots Isolation Forest Outliers Identification Grubbs' Test Z-Score Standard Deviation Method Interquartile Range Method Tukey's method Internally studentized residuals AKA z-score method Median Absolute Deviation method
	Outliers - MultiVariate
	Visualizing Outliers ScatterPlots Outliers Identification Mahalanobis Distance Robust Mahalanobis Distance DBSCAN Clustering PyOD Library
	Automatic Detection and Removal of Outliers
	Compare Algorithms LocalOutlierFactor IsolationForest Minimum Covariance Determinant

0.2 Consideraciones casos CitizenLab programados en R

- La mayoría de las tareas de este proceso se han realizado en los notebooks del proceso 05 Data Collection porque eran necesarias para las tareas ETL. En esos casos, en este notebook se referencia al notebook del proceso 05 correspondiente
- Por tanto en los notebooks de este proceso de manera general se incluyen las comprobaciones necesarias, y comentarios si procede
- Las tareas del proceso se van a aplicar solo a los archivos que forman parte del despliegue, ya que hay muchos archivos intermedios que no procede pasar por este proceso
- El nombre de archivo del notebook hace referencia al nombre de archivo del proceso 05 al que se aplica este proceso, por eso pueden no ser correlativa la numeración
- Las comprobaciones se van a realizar teniendo en cuenta que el lenguaje utilizado en el despliegue de este caso es R

0.3 File

- Input File: CU_18_09.2_20_diario_infra
- Output File: CU_18_09.3_20_diario_infra

0.3.1 Encoding

Con la siguiente expresión se evitan problemas con el encoding al ejecutar el notebook. Es posible que deba ser eliminada o adaptada a la máquina en la que se ejecute el código.

```
[1]: Sys.setlocale(category = "LC_ALL", locale = "es_ES.UTF-8")
```

```
'LC_CTYPE=es_ES.UTF-8;LC_NUMERIC=C;LC_TIME=es_ES.UTF-8;LC_COLLATE=es_ES.UTF-8;LC_MONETARY=es_ES.UTF-8;LC_MESSAGES=en_US.UTF-8;LC_PAPER=es_ES.UTF-8;LC_NAME=C;LC_ADDRESS=C;LC_TELEPHONE=C;LC_MEASUREMENT=es_ES.UTF-8;LC_IDENTIFICATION=C'
```

0.4 Settings

0.4.1 Libraries to use

```
[2]: library(readr)
library(dplyr)
# library(sf)
library(tidyr)
library(stringr)
library(ggplot2)
```

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

```
intersect, setdiff, setequal, union
```

0.4.2 Paths

```
[3]: iPath <- "Data/Input/"
     oPath <- "Data/Output/"
```

0.5 Data Load

OPCION A: Seleccionar fichero en ventana para mayor comodidad

Data load using the {tcltk} package. Uncomment the line if using this option

```
[4]: # file_data <- tcltk::tk_choose.files(multi = FALSE)
```

OPCION B: Especificar el nombre de archivo

```
[5]: iFile <- "CU_18_09.2_20_diario_infra.csv"
     file_data <- paste0(iPath, iFile)

     if(file.exists(file_data)){
       cat("Se leerán datos del archivo: ", file_data)
     } else{
       warning("Cuidado: el archivo no existe.")
     }
}
```

Se leerán datos del archivo: Data/Input/CU_18_09.2_20_diario_infra.csv

Data file to dataframe Usar la función adecuada según el formato de entrada (xlsx, csv, json, ...)

```
[6]: data <- read_csv(file_data)
```

Rows: 377816 Columns: 10

Column specification

Delimiter: ","

dbl (9): id_inf, capacidad, demanda, evento_infra, evento_zona, tmed,
prec,...

date (1): fecha

Use `spec()` to retrieve the full column specification for this data.

Specify the column types or set `show_col_types = FALSE` to quiet this message.

Visualizo los datos.

Estructura de los datos:

```
[7]: data |> glimpse()
```

```
Rows: 377,816
Columns: 10
$ id_inf      <dbl> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12,
13, 14, 15, 16, 17...
$ fecha       <date> 2019-01-01, 2019-01-01, 2019-01-01,
2019-01-01, 2019-01-...
$ capacidad   <dbl> 993, 996, 1036, 1020, 992, 1026, 1007,
976, 1037, 972, 94...
$ demanda     <dbl> 883, 888, 922, 1134, 1103, 1139, 897,
1086, 1150, 861, 83...
$ evento_infra <dbl> 1, 0, 0, 1, 1, 0, 0, 1, 1, 0, 1, 1, 0,
0, 0, 1, 0, 1, 1, ...
$ evento_zona <dbl> 1, 0, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0, 0,
1, 1, 1, 1, 0, 1, ...
$ tmed        <dbl> 6.953211, 6.196420, 6.483569, 5.875797,
6.212680, 5.87854...
$ prec        <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, ...
$ velmedia    <dbl> 0.6433886, 0.3417523, 0.4132169,
0.1820178, 0.2118110, 0.1...
$ presMax     <dbl> 952.9357, 950.2191, 950.8051, 951.6768,
953.5118, 952.168...
```

Muestra de los primeros datos:

```
[8]: data |> slice_head(n = 5)
```

	id_inf	fecha	capacidad	demanda	evento_infra	evento_zona	tmed	p
	<dbl>	<date>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
	1	2019-01-01	993	883	1	1	6.953211	0
A spec_tbl_df: 5 × 10	2	2019-01-01	996	888	0	0	6.196420	0
	3	2019-01-01	1036	922	0	0	6.483569	0
	4	2019-01-01	1020	1134	1	0	5.875797	0
	5	2019-01-01	992	1103	1	1	6.212680	0

0.6 Outliers - Univariate

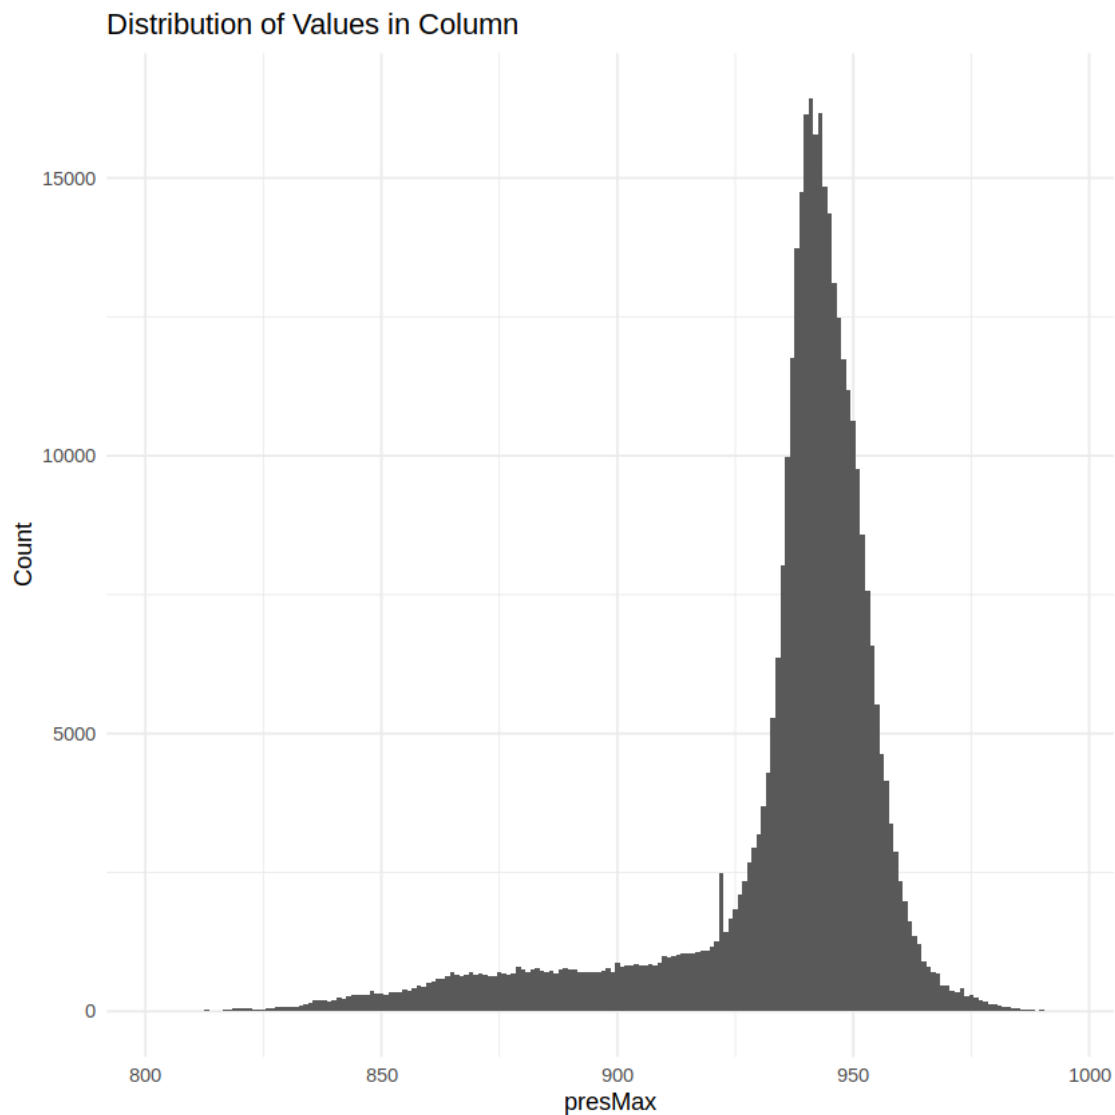
0.6.1 Visualizing Outliers

Distribution Selecting feature to analyze

```
[9]: # Selecting feature to analyze
column_name <- "presMax"
```

Operation

```
[10]: # Create a histogram plot
ggplot(data, aes(x = data[[column_name]])) +
  geom_histogram(binwidth = 1) +
  labs(x = column_name, y = "Count") +
  ggtitle("Distribution of Values in Column") +
  theme_minimal()
```



Box Plots Are great to summarize and visualize the distribution of variables easily and quickly.


```
[ ]:
```

Selecting feature to analyze

```
[11]: # Selecting feature to analyze  
column_name <- "presMax"
```

Operation

```
[12]: # Analyze specifics features  
# Create a box plot  
ggplot(data, aes(x = "", y = data[[column_name]])) +  
  geom_boxplot() +  
  labs(x = "", y = column_name) +  
  ggtitle("Box Plot of Column") +  
  theme_minimal()
```



Isolation Forest Selecting feature to analyze

```
[13]: # Selecting feature to analyze
```

Operation

```
[ ]:
```

0.6.2 Outliers Identification

Grubbs' Test Selecting feature to analyze

```
[14]: # Selecting feature to analyze  
column_name <- "presMax"
```

Operation

```
[15]: library(outliers)  
outliers <- grubbs.test(data[[column_name]], opposite = FALSE)  
print(outliers)
```

Grubbs test for one outlier

```
data: data[[column_name]]  
G = 5.48406, U = 0.99992, p-value = 0.007848  
alternative hypothesis: lowest value 801.911968064451 is an outlier
```

Z-Score Selecting feature to analyze

```
[16]: # Selecting feature to analyze  
column_name <- "presMax"
```

Operation

```
[ ]:
```

```
[18]: # Define a threshold to identify an outlier.  
# List of row numbers with outlier  
# Choose the numeric column from your data  
  
# Calculate the z-score  
z_scores <- data %>% select(all_of(column_name)) %>% scale()
```

```

# Define a threshold for identifying outliers (e.g., z-score > 5 or z-score <
↪-5)
threshold <- 5

# Find the row numbers with z-scores exceeding the threshold
outlier_rows <- which(abs(z_scores) > threshold)

# Print the row numbers with outliers
print(outlier_rows)

data_cleaned <- subset(data, !(row.names(data) %in% outlier_rows))

```

```

[1] 17176 18314 18315 19452 19453 20590 20591 22866 24004 24005
[11] 24682 25142 25820 28556 28557 29694 29695 145770 145771 146908
[21] 146909 165116 165117 166254 166255 171944 171945 173082 173083 174220
[31] 174221 305090 310780 310781 311918 311919 314194 314195 327850 327851
[41] 330126 330127 333540 333541 334678 334679 335816 335817 338092 338093
[51] 338770 339230 339231 340368 340369 341506 341507 342644 342645 343322
[61] 343782 344920 344921 346058 346059 348334 348335 349012 349472 350150
[71] 350610 350611 358576 358577 359714 359715 360392 360852 360853 361990
[81] 361991 363128 363129 364266 364267 364944 365404 365405 366082

```

Standard Deviation Method Selecting feature to analyze

```

[38]: # Selecting all features to analyze
column_name <- "presMax"

```

Operation

```

[ ]:

```

```

[44]: ## identify outliers with standard deviation
## Choose the numeric column from your data
# column_name_df <- data[[column_name]]

## Calculate the mean and standard deviation of the column
# column_mean <- mean(column_name_df)
# column_sd <- sd(column_name_df)

## Define the threshold as a multiple of the standard deviation (e.g., 3 times
↪the standard deviation)
# threshold <- 3

## Identify the outliers based on the threshold
# outlier_rows <- which(column_name_df > (column_mean + threshold * column_sd)
↪| column_name_df < (column_mean - threshold * column_sd))

```

```
# # Print the updated column with outliers removed
# print(outlier_rows)

# data_cleaned <- subset(data, !(row.names(data) %in% outlier_rows))
```

Interquartile Range Method Selecting factor k

```
[45]: # Selecting factor k
column_name <- "presMax"
threshold <- 1.5
```

Operation

```
[48]: # identify outliers with standard deviation
# Choose the numeric column from your data
# column_name_df <- data[[column_name]]

# # Calculate the first quartile (Q1) and third quartile (Q3)
# Q1 <- quantile(column_name_df, 0.25)
# Q3 <- quantile(column_name_df, 0.75)

# # Calculate the IQR (Interquartile Range)
# IQR <- Q3 - Q1

# # Identify the outliers based on the threshold
# outlier_rows <- which(column_name_df < (Q1 - threshold * IQR) |
  ↪ column_name_df > (Q3 + threshold * IQR))
# print(outlier_rows)
# data_cleaned <- subset(data, !(row.names(data) %in% outlier_rows))
```

Tukey's method

```
[51]: #Tukey's method
# column_name <- "presMax"

# column_name_df <- data[[column_name]]
# # Calculate the first quartile (Q1) and third quartile (Q3)
# Q1 <- quantile(column_name_df, 0.25)
# Q3 <- quantile(column_name_df, 0.75)

# # Calculate the interquartile range (IQR)
# IQR <- Q3 - Q1

# # Define the multiplier for Tukey's method (e.g., 1.5 times the IQR)
# multiplier <- 1.5

# # Calculate the lower and upper bounds for outliers
```

```
# lower_bound <- Q1 - multiplier * IQR
# upper_bound <- Q3 + multiplier * IQR

# # Identify the outliers based on the bounds
# outlier_rows <- which(column_name_df < lower_bound | column_name_df >
  ↪upper_bound)

# # Print the updated column with outliers removed
# print(outlier_rows)
# data_cleaned <- subset(data, !(row.names(data) %in% outlier_rows))
```

Internally studentized residuals AKA z-score method

```
[ ]: #Internally studentized method (z-score)
```

Median Absolute Deviation method

```
[53]: #MAD method
# column_name <- "presMax"

# column_name_df <- data[[column_name]]
# # Calculate the median absolute deviation (MAD)
# mad <- median(abs(column_name_df - median(column_name_df, na.rm = TRUE)), na.
  ↪rm = TRUE)

# # Define a threshold for identifying outliers (e.g., 3 times the MAD)
# threshold <- 3 * mad

# # Identify the outliers based on the MAD
# outliers <- which(abs(column_name_df - median(column_name_df, na.rm = TRUE))
  ↪threshold)

# # Print the updated column with outliers removed
# print(outliers)
# data_cleaned <- subset(data, !(row.names(data) %in% outliers))
```

0.7 Outliers - MultiVariate

0.7.1 Visualizing Outliers

ScatterPlots: a common way to plot multivariate outliers is the scatter plot.

ScatterPlots A common way to plot multivariate outliers is the scatter plot.

```
[ ]:
```

Selecting feature to analyze

```
[ ]: # Selecting feature to analyze
```

Operation

```
[ ]:
```

0.7.2 Outliers Identification

Mahalanobis Distance

```
[ ]:
```

Selecting feature to analyze

```
[ ]: # Selecting features to analyze
```

Operation

```
[ ]: # Analyze selected features
```

```
[ ]: # Analyze all dataset
```

Robust Mahalanobis Distance

```
[ ]: #Robust Mahalonibis Distance
```

Selecting feature to analyze

```
[ ]: # Selecting features to analyze
```

Operation

```
[ ]: # Analyze selected features
```

```
[ ]: # Analyze all dataset
```

DBSCAN Clustering Selecting feature to analyze

```
[ ]: # Selecting feature to analyze
```

Operation

```
[ ]:
```

```
[ ]: # specify & fit model
```

```
[ ]: # visualize outputs
```

```
[ ]: # outliers dataframe
```

```
[ ]: # Index of rows with outliers
```

```
[ ]: # Outliers Dataframe
```

```
[ ]:
```

0.8 Data Save

- Solo si se han hecho cambios
- No aplica

Identificamos los datos a guardar

```
[22]: data_to_save <- data_cleaned
```

Estructura de nombre de archivos:

- Código del caso de uso, por ejemplo "CU_04"
- Número del proceso que lo genera, por ejemplo "_06".
- Resto del nombre del archivo de entrada
- Extensión del archivo

Ejemplo: "CU_04_06_01_01_zonasgeo.json, primer fichero que se genera en la tarea 01 del proceso 05 (Data Collection) para el caso de uso 04 (vacunas) y que se ha transformado en el proceso 06

Importante mantener los guiones bajos antes de proceso, tarea, archivo y nombre

0.8.1 Proceso 09.2

```
[21]: caso <- "CU_18"  
      proceso <- '_09.3'  
      tarea <- "_20"  
      archivo <- ""  
      proper <- "_diario_infra"  
      extension <- ".csv"
```

OPCION A: Uso del paquete "tcltk" para mayor comodidad

- Buscar carpeta, escribir nombre de archivo SIN extensión (se especifica en el código)
- Especificar sufijo2 si es necesario
- Cambiar datos por datos_xx si es necesario

```
[ ]: # file_save <- paste0(caso, proceso, tarea, tcltk::tkgetSaveFile(), proper,  
    ↪ extension)  
# path_out <- paste0(oPath, file_save)  
# write_csv(data_to_save_XXXXX, path_out)  
  
# cat('File saved as: ')  
# path_out
```

OPCION B: Especificar el nombre de archivo

- Los ficheros de salida del proceso van siempre a Data/Output/.

```
[23]: file_save <- paste0(caso, proceso, tarea, archivo, proper, extension)
      path_out <- paste0(oPath, file_save)
      write_csv(data_to_save, path_out)

      cat('File saved as: ')
      path_out
```

File saved as:

'Data/Output/CU_18_09.3_20_diario_infra.csv'

Copia del fichero a Input Si el archivo se va a usar en otros notebooks, copiar a la carpeta Input

```
[24]: path_in <- paste0(iPath, file_save)
      file.copy(path_out, path_in, overwrite = TRUE)
```

TRUE

0.9 REPORT

A continuación se realizará un informe de las acciones realizadas

0.10 Main Actions Carried Out

- Si eran necesarias se han realizado en el proceso 05 por cuestiones de eficiencia

0.11 Main Conclusions

- Los datos están limpios para el despliegue

0.12 CODE TO DEPLOY (PILOT)

A continuación se incluirá el código que deba ser llevado a despliegue para producción, dado que se entiende efectúa operaciones necesarias sobre los datos en la ejecución del prototipo

Description

- No hay nada que desplegar en el piloto, ya que estos datos son estáticos o en todo caso cambian con muy poca frecuencia, altamente improbable durante el proyecto.

CODE

```
[ ]:
```